

Bayesian Inference for Deep Learning

Inference and modern trends for Bayesian Neural Networks: Variational Inference

Simone Rossi and Maurizio Filippone

Data Science Department, EURECOM (France)

Roadmap of the tutorial

- **Introduction to Bayesian Inference**
- **Bayesian inference as Optimization with Variational Inference**
 - Introduction to variational inference (objective and gradients)
 - Challenges and solutions for variational inference on Bayesian neural networks
- **Sampling with MCMC methods**
- **Alternatives for Approximate Bayesian Deep Learning**
- **Gaussian processes and Bayesian neural networks**
- **Priors and Model Selection**
- **Uncertainty Quantification with Bayesian Neural Networks**

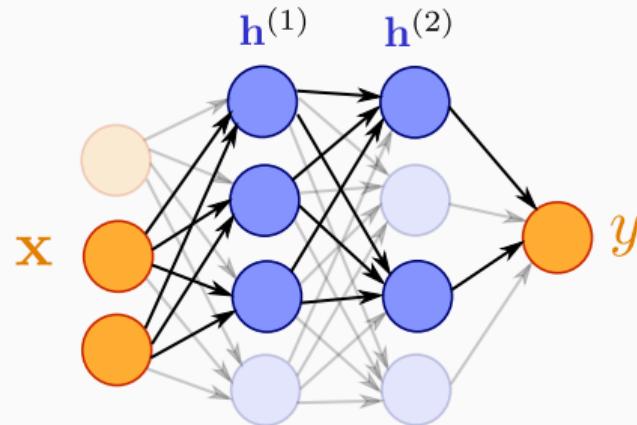
An Introduction to Bayesian Neural Networks with Monte-Carlo Dropout

Let's reuse what we know

Dropout is a simple and powerful method to avoid overfitting in deep neural networks.

Let's reuse what we know

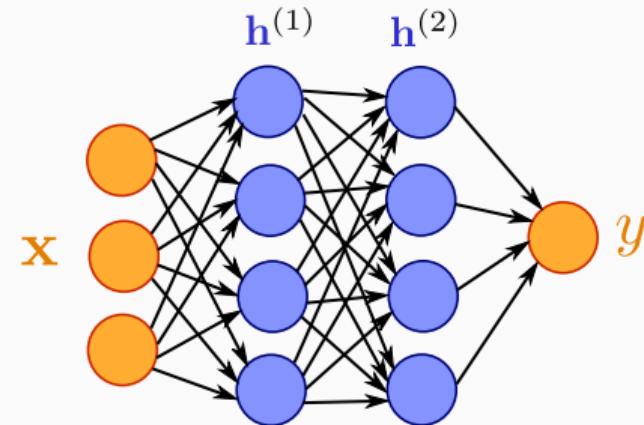
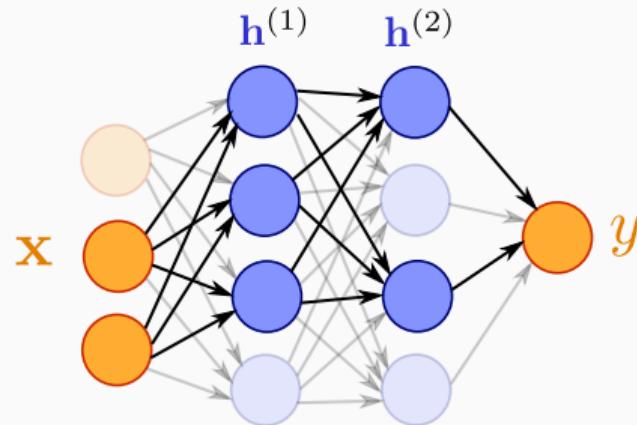
Dropout is a simple and powerful method to avoid overfitting in deep neural networks.



At each iteration of **train time**, some units are dropped (with probability p).

Let's reuse what we know

Dropout is a simple and powerful method to avoid overfitting in deep neural networks.

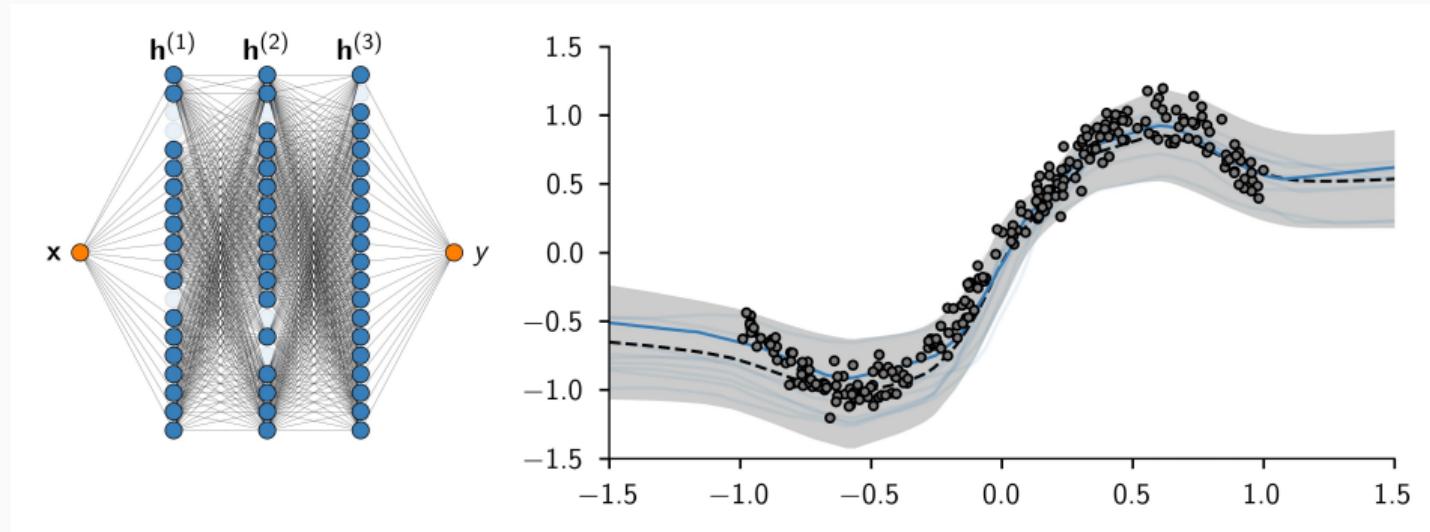


At each iteration of **train time**, some units are dropped (with probability p).

At **test time**, all units are considered.

Srivastava et al. (2014). *Dropout: A Simple Way to Prevent Neural Networks from Overfitting*, JMLR

What happens if we use multiple dropout masks also at test time?



Compute the mean prediction as $\hat{y}_* \approx \frac{1}{T} \sum_t^T f(x_*; [\{\tilde{W}^{(1)}, \dots, \tilde{W}^{(L)}\}]_t)$ and evaluate the variance at test point $\text{Var}[\hat{y}_*]$

A bit more formal

- Training with **regularized loss** ...

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N e(\mathbf{y}_i, f(\mathbf{x}_i)) + \lambda \sum_{l=1}^L \|\mathbf{W}^{(l)}\|^2$$

... equivalent to do approximate Bayesian inference.

A bit more formal

- Training with **regularized loss** ...

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N e(\mathbf{y}_i, f(\mathbf{x}_i)) + \lambda \sum_{l=1}^L \|\mathbf{W}^{(l)}\|^2$$

... equivalent to do approximate Bayesian inference.

- If we define $\mathbf{w} = \{\mathbf{W}_1, \dots, \mathbf{W}_L\}$ and a “posterior” distribution $q(\mathbf{w})$ such that,

$$\mathbf{W}_i = \mathbf{M}_i \cdot \text{diag}(\mathbf{z}_i), \text{ with } z_{ij} \sim \text{Bern}(p_i)$$

we can use the following objective

$$\mathcal{L}_{\text{ELBO}} = - \sum_{i=1}^N \int \log p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w}) q(\mathbf{w}) d\mathbf{w} + \text{KL}[q(\mathbf{w}) \parallel p(\mathbf{w})]$$

Why? What is this?

An Introduction to Variational Inference

Our problem setup

- Model: $f(\mathbf{x}; \mathbf{w})$ is a L-layer neural network with weights $\mathbf{w} = \{\mathbf{W}^{(1)}, \dots, \mathbf{W}^{(L)}\}$
- Likelihood:

$$p(\mathbf{Y} | \mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(\mathbf{y}_n | \mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N p(\mathbf{y}_n | f(\mathbf{x}_n; \mathbf{w}))$$

- Prior:
- $$p(\mathbf{w}) = \prod_{l=1}^L p(\mathbf{W}^{(l)}) = \prod_{i=1}^L \prod_{jk} \mathcal{N}(W_{jk}^{(l)} | 0, 1)$$
- Posterior (using Bayes's theorem):

$$p(\mathbf{w} | \mathbf{Y}, \mathbf{X}) = \frac{p(\mathbf{Y} | \mathbf{X}, \mathbf{w}) p(\mathbf{w})}{p(\mathbf{Y} | \mathbf{X})}$$

Our problem setup

- Model: $f(\mathbf{x}; \mathbf{w})$ is a L-layer neural network with weights $\mathbf{w} = \{\mathbf{W}^{(1)}, \dots, \mathbf{W}^{(L)}\}$
- Likelihood:

$$p(\mathbf{Y} | \mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(\mathbf{y}_n | \mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N p(\mathbf{y}_n | f(\mathbf{x}_n; \mathbf{w}))$$

- Prior:
- $$p(\mathbf{w}) = \prod_{l=1}^L p(\mathbf{W}^{(l)}) = \prod_{i=1}^L \prod_{jk} \mathcal{N}(W_{jk}^{(l)} | 0, 1)$$
- Posterior (using Bayes's theorem):

$$p(\mathbf{w} | \mathbf{Y}, \mathbf{X}) = \frac{p(\mathbf{Y} | \mathbf{X}, \mathbf{w}) p(\mathbf{w})}{p(\mathbf{Y} | \mathbf{X})}$$

We can't compute the marginal likelihood ...

$$p(\mathbf{Y} | \mathbf{X}) = \int p(\mathbf{Y} | \mathbf{X}, \mathbf{w}) p(\mathbf{w}) d\mathbf{w}$$

... and we can't compute the predictive distribution:

$$p(\mathbf{y}_* | \mathbf{x}_*, \mathbf{y}, \mathbf{X}) = \int p(\mathbf{y}_* | \mathbf{x}_*, \mathbf{w}) p(\mathbf{w} | \mathbf{Y}, \mathbf{X}) d\mathbf{w}$$

Variational inference

Solutions:

- Collapse the posterior on the most likely value (Maximum-a-Posteriori or MAP)
- Approximate the intractable posterior:
 - Use Variational Inference
- Local approximation:
 - Use Laplace approximation (local approximation at the MAP solution)
- Sample from the intractable posterior:
 - Markov-Chain Monte-Carlo (Hamiltonian Monte-Carlo)

Variational inference

Solutions:

- Collapse the posterior on the most likely value (Maximum-a-Posteriori or MAP)
- Approximate the intractable posterior:
 - Use Variational Inference
- Local approximation:
 - Use Laplace approximation (local approximation at the MAP solution)
- Sample from the intractable posterior:
 - Markov-Chain Monte-Carlo (Hamiltonian Monte-Carlo)

Variational Inference

Main idea

Instead of trying to solve intractable integrals, let's solve an optimization problem.

Main idea

Instead of trying to solve intractable integrals, let's solve an optimization problem.

A very general recipe:

- Introduce a set \mathcal{Q} of distributions $q(\mathbf{w})$
- Define an objective which measures the “distance” between an arbitrary distribution $q(\mathbf{w}) \in \mathcal{Q}$ and $p(\mathbf{w} | \mathbf{Y}, \mathbf{X})$
- In the set of possible solutions \mathcal{Q} , find the best $q(\mathbf{w})$ that minimizes the “distance” to $p(\mathbf{w} | \mathbf{Y}, \mathbf{X})$

Interpret $q(\mathbf{w})$ as a distribution that approximates the intractable $p(\mathbf{w} | \mathbf{Y}, \mathbf{X})$

Visual illustration of Variational Inference

Space of all possible solutions given a likelihood/prior pair

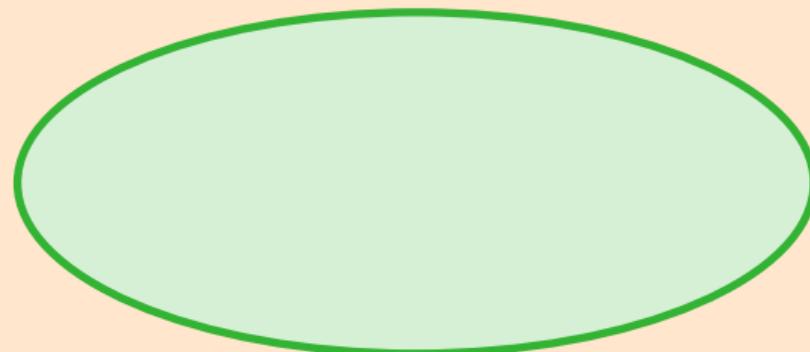
$$p(\text{w} \mid \text{Y}, \text{X})$$

Visual illustration of Variational Inference

Space of all possible solutions given a likelihood/prior pair

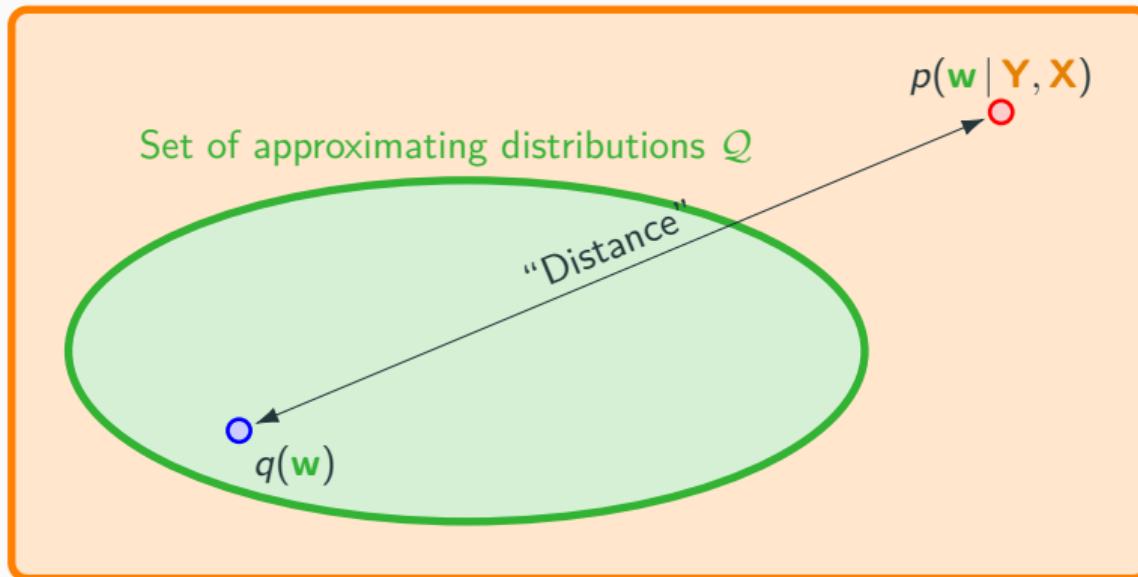
$$p(\mathbf{w} \mid \mathbf{Y}, \mathbf{X})$$

Set of approximating distributions \mathcal{Q}



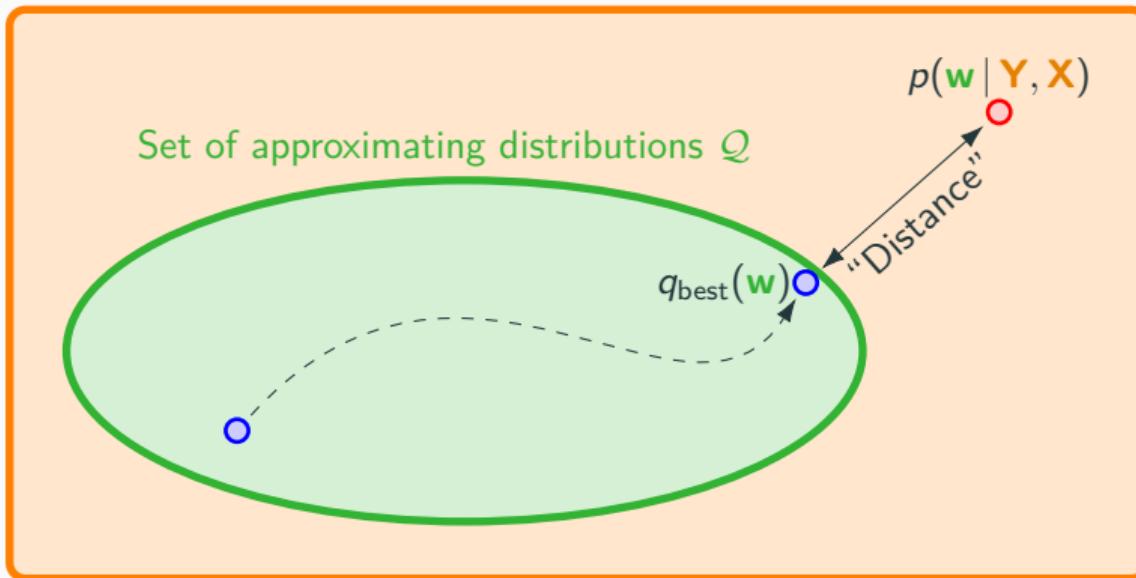
Visual illustration of Variational Inference

Space of all possible solutions given a likelihood/prior pair



Visual illustration of Variational Inference

Space of all possible solutions given a likelihood/prior pair



Variational Inference - Form of the approximation

- The *mean-field approach* imposes independent distributions for each component of \mathbf{w}

$$q(\mathbf{w}) = \prod_{l=1}^L \prod_{i=1}^{D_{in}^{(l)}} \prod_{j=1}^{D_{out}^{(l)}} q(W_{ij}^{(l)})$$

For notation purposes, all weights are concatenated in $\mathbf{w} \in \mathbb{R}^P$ (P is the number of weights in the network)

Variational Inference - Form of the approximation

- The *mean-field approach* imposes independent distributions for each component of \mathbf{w}

$$q(\mathbf{w}) = \prod_{l=1}^L \prod_{i=1}^{D_{in}^{(l)}} \prod_{j=1}^{D_{out}^{(l)}} q(\mathbf{W}_{ij}^{(l)})$$

For notation purposes, all weights are concatenated in $\mathbf{w} \in \mathbb{R}^P$ (P is the number of weights in the network)

- We start with Gaussian distributions

$$q(\mathbf{w}) = \prod_{i=1}^P q(\mathbf{w}_i) = \prod_{i=1}^P \mathcal{N}(\mathbf{w}_i | \boldsymbol{\mu}_i, \sigma_i^2)$$

$\boldsymbol{\mu}_i, \sigma_i^2$ are called *variational parameters* (for notation, they are collected into $\boldsymbol{\nu}$)

Variational Inference - Objective

- We will use the KL divergence to measure the “distance” between the two distributions,

$$\text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w}|\mathbf{Y}, \mathbf{X})] = \int q(\mathbf{w}; \boldsymbol{\nu}) \log \frac{q(\mathbf{w}; \boldsymbol{\nu})}{p(\mathbf{w}|\mathbf{Y}, \mathbf{X})} d\mathbf{w} = E_{q(\mathbf{w}; \boldsymbol{\nu})} \left[\log \frac{q(\mathbf{w}; \boldsymbol{\nu})}{p(\mathbf{w}|\mathbf{Y}, \mathbf{X})} \right]$$

- A tractable objective to optimize $q(\mathbf{w}; \boldsymbol{\nu})$ is obtained by manipulating the KL divergence

$$\text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w}|\mathbf{Y}, \mathbf{X})] = -E_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] + \text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w})] + \log p(\mathbf{Y}|\mathbf{X})$$

Variational Inference - Objective

- We will use the KL divergence to measure the “distance” between the two distributions,

$$\text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w}|\mathbf{Y}, \mathbf{X})] = \int q(\mathbf{w}; \boldsymbol{\nu}) \log \frac{q(\mathbf{w}; \boldsymbol{\nu})}{p(\mathbf{w}|\mathbf{Y}, \mathbf{X})} d\mathbf{w} = E_{q(\mathbf{w}; \boldsymbol{\nu})} \left[\log \frac{q(\mathbf{w}; \boldsymbol{\nu})}{p(\mathbf{w}|\mathbf{Y}, \mathbf{X})} \right]$$

- A tractable objective to optimize $q(\mathbf{w}; \boldsymbol{\nu})$ is obtained by manipulating the KL divergence

$$\text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w}|\mathbf{Y}, \mathbf{X})] = -E_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] + \text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w})] + \log p(\mathbf{Y}|\mathbf{X})$$

- The last term is the problematic one ...

Variational Inference - Objective

- ... but manipulating the previous expression

$$\log p(\mathbf{Y}|\mathbf{X}) - \text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w}|\mathbf{Y}, \mathbf{X})] = \underbrace{\mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] - \text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w})]}_{\mathcal{L}_{\text{ELBO}}}$$

we have a computable objective

- Minimizing the original KL is equivalent to maximize $\mathcal{L}_{\text{ELBO}}$

Variational Inference - Objective

- ... but manipulating the previous expression

$$\log p(\mathbf{Y}|\mathbf{X}) - \text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w}|\mathbf{Y}, \mathbf{X})] = \underbrace{\mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] - \text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w})]}_{\mathcal{L}_{\text{ELBO}}}$$

we have a computable objective

- Minimizing the original KL is equivalent to maximize $\mathcal{L}_{\text{ELBO}}$

Note that $\text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \parallel p(\mathbf{w}|\mathbf{Y}, \mathbf{X})] \geq 0$:

- The right hand side is a lower bound to the marginal likelihood (hence the name)
- If we can make $q(\mathbf{w}; \boldsymbol{\nu})$ equal to the posterior, our objective will be equal to the marginal likelihood!

Variational Inference - Objective

Variational objective (a.k.a. evidence lower bound) to be maximized wrt $q(\mathbf{w}; \boldsymbol{\nu})$

$$\mathcal{L}_{\text{ELBO}} = \mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] - \text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \| p(\mathbf{w})]$$

Variational Inference - Objective

Variational objective (a.k.a. evidence lower bound) to be maximized wrt $q(\mathbf{w}; \boldsymbol{\nu})$

$$\mathcal{L}_{\text{ELBO}} = E_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] - \text{KL} [q(\mathbf{w}; \boldsymbol{\nu}) \| p(\mathbf{w})]$$

- First term is a model fitting term:

$$E_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})]$$

the higher the better the parameters drawn from $q(\mathbf{w}; \boldsymbol{\nu})$ are at modeling the labels

Variational Inference - Objective

Variational objective (a.k.a. evidence lower bound) to be maximized wrt $q(\mathbf{w}; \boldsymbol{\nu})$

$$\mathcal{L}_{\text{ELBO}} = E_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] - \text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \| p(\mathbf{w})]$$

- First term is a model fitting term:

$$E_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})]$$

the higher the better the parameters drawn from $q(\mathbf{w}; \boldsymbol{\nu})$ are at modeling the labels

- Second term is a regularization term:

$$-\text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \| p(\mathbf{w})]$$

which penalizes $q(\mathbf{w}; \boldsymbol{\nu})$ deviating too much from the prior

Variational Inference - Computation

$$\mathcal{L}_{\text{ELBO}} = \mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y} | \mathbf{X}, \mathbf{w})] - \text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \| p(\mathbf{w})]$$

- The second term can be expressed analytically by using the expression of the KL divergence and it does **not** depend on the neural network architecture

Variational Inference - Computation

$$\mathcal{L}_{\text{ELBO}} = \mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] - \text{KL}[q(\mathbf{w}; \boldsymbol{\nu}) \| p(\mathbf{w})]$$

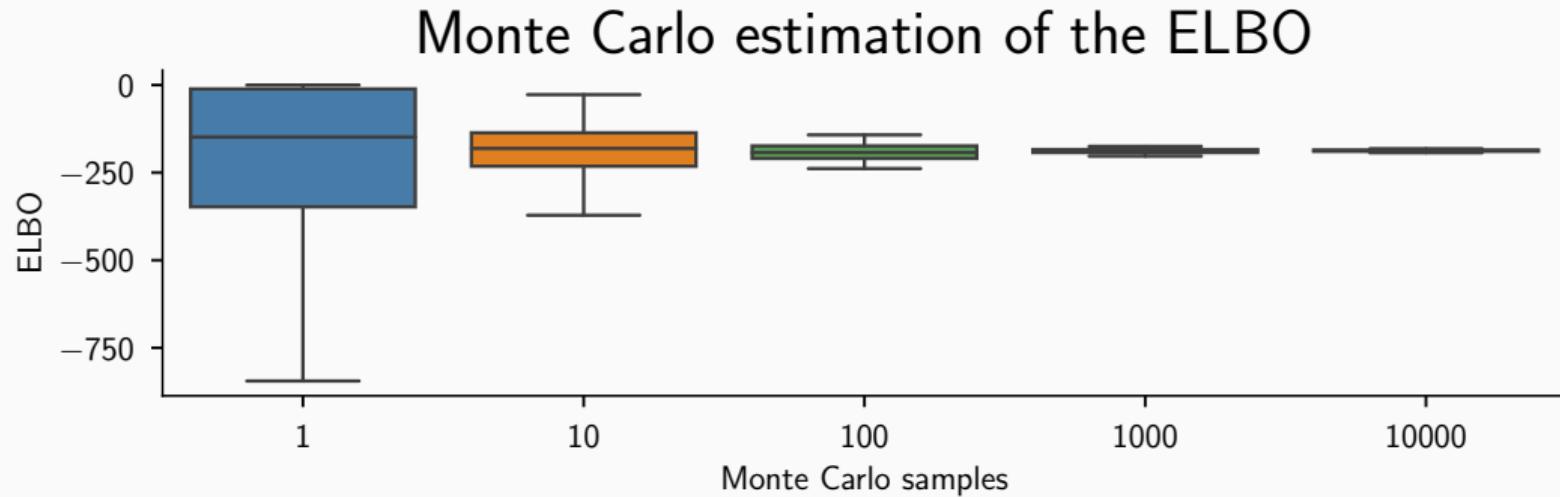
- The second term can be expressed analytically by using the expression of the KL divergence and it does **not** depend on the neural network architecture
- The first expectation is never analytically available for Bayesian neural networks but it can be approximated using Monte Carlo integration:

$$\mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] \approx \frac{1}{N_{\text{MC}}} \sum_{i=1}^{N_{\text{MC}}} \log p(\mathbf{Y}|\mathbf{X}, \tilde{\mathbf{w}}^{(h)}), \quad \tilde{\mathbf{w}}^{(h)} \sim q(\mathbf{w}; \boldsymbol{\nu})$$

Variational Inference - Computation

$$\mathcal{L}_{\text{ELBO}} = \frac{1}{N_{\text{MC}}} \sum_{i=1}^{N_{\text{MC}}} \log p(\mathbf{Y}|\mathbf{X}, \tilde{\mathbf{w}}^{(h)}) - \text{KL}[q(\mathbf{w}; \nu) \| p(\mathbf{w})]$$

Remember: this estimator is unbiased and its variance shrinks $\propto 1/N_{\text{MC}}$, independently of the dimensionality.



Variational Inference - Optimization

We need to compute the gradients of the objective w.r.t the *variational parameters* ν

$$\nabla_{\nu} \mathcal{L}_{\text{ELBO}} = \nabla_{\nu} \mathbb{E}_{q(\mathbf{w}; \nu)} [\log p(\mathbf{Y} | \mathbf{X}, \mathbf{w})] - \nabla_{\nu} \text{KL}[q(\mathbf{w}; \nu) \| p(\mathbf{w})]$$

- The second term is easy (everything is deterministic)

Variational Inference - Optimization

We need to compute the gradients of the objective w.r.t the *variational parameters* ν

$$\nabla_{\nu} \mathcal{L}_{\text{ELBO}} = \nabla_{\nu} \mathbb{E}_{q(\mathbf{w}; \nu)} [\log p(\mathbf{Y} | \mathbf{X}, \mathbf{w})] - \nabla_{\nu} \text{KL}[q(\mathbf{w}; \nu) \| p(\mathbf{w})]$$

- The second term is easy (everything is deterministic)
- The first term requires a bit of work

Variational Inference - Optimization

- The Monte Carlo approximation

$$\mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] \approx \frac{1}{N_{MC}} \sum_{h=1}^{N_{MC}} \log p(\mathbf{Y}|\mathbf{X}, \tilde{\mathbf{w}}^{(h)})$$

makes it tricky to optimize the objective

- With $q(\mathbf{w}; \boldsymbol{\nu})$ fixed, when we resample \mathbf{w} from $q(\mathbf{w}; \boldsymbol{\nu})$ we obtain a different value!
- How can we make gradient updates to the μ_i, σ_i^2 parameters of $q(\mathbf{w}; \boldsymbol{\nu})$?

Variational Inference - Optimization

- The Monte Carlo approximation

$$E_{q(\mathbf{w}; \boldsymbol{\nu})} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] \approx \frac{1}{N_{MC}} \sum_{h=1}^{N_{MC}} \log p(\mathbf{Y}|\mathbf{X}, \tilde{\mathbf{w}}^{(h)})$$

makes it tricky to optimize the objective

- With $q(\mathbf{w}; \boldsymbol{\nu})$ fixed, when we resample \mathbf{w} from $q(\mathbf{w}; \boldsymbol{\nu})$ we obtain a different value!
- How can we make gradient updates to the μ_i, σ_i^2 parameters of $q(\mathbf{w}; \boldsymbol{\nu})$?
- Answer: freeze the randomness within Monte Carlo!

Variational Inference - Reparameterization trick

$$\nabla_{\nu} E_{q(w; \nu)} \log p(Y|X, w)$$

Idea:

- Samples of w can be obtained by a *deterministic transformation* ϕ of a random variable $\varepsilon \sim p(\varepsilon)$, such that $p(\varepsilon)$ does not depend on the variational parameters
- The variational parameters ν are parameters of the function ϕ
- Use the chain rule of differentiation to push the gradient through this function ϕ



Variational Inference - Reparameterization trick (Derivation)

Key observation

$$\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) = E_{p(\varepsilon)} \nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}$$

Now, we can turn the *gradient of an expectation* into an *expectation of a gradient*.

$$\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) = \nabla_{\nu} E_{p(\varepsilon)} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}$$

Variational Inference - Reparameterization trick (Derivation)

Key observation

$$\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) = E_{p(\varepsilon)} \nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}$$

Now, we can turn the *gradient of an expectation* into an *expectation of a gradient*.

$$\begin{aligned}\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) &= \nabla_{\nu} E_{p(\varepsilon)} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \\ &= E_{p(\varepsilon)} [\nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}]\end{aligned}$$

Variational Inference - Reparameterization trick (Derivation)

Key observation

$$\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) = E_{p(\varepsilon)} \nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}$$

Now, we can turn the *gradient of an expectation* into an *expectation of a gradient*.

$$\begin{aligned}\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) &= \nabla_{\nu} E_{p(\varepsilon)} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \\ &= E_{p(\varepsilon)} [\nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}] \\ &= E_{p(\varepsilon)} [\nabla_w \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \nabla_{\nu} w]\end{aligned}$$

Variational Inference - Reparameterization trick (Derivation)

Key observation

$$\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) = E_{p(\varepsilon)} \nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}$$

Now, we can turn the *gradient of an expectation* into an *expectation of a gradient*.

$$\begin{aligned}\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) &= \nabla_{\nu} E_{p(\varepsilon)} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \\ &= E_{p(\varepsilon)} [\nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}] \\ &= E_{p(\varepsilon)} [\nabla_w \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \nabla_{\nu} \phi(\varepsilon; \nu)]\end{aligned}$$

Variational Inference - Reparameterization trick (Derivation)

Key observation

$$\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) = E_{p(\varepsilon)} \nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}$$

Now, we can turn the *gradient of an expectation* into an *expectation of a gradient*.

$$\begin{aligned}\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) &= \nabla_{\nu} E_{p(\varepsilon)} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \\ &= E_{p(\varepsilon)} [\nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}] \\ &= E_{p(\varepsilon)} [\nabla_w \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \nabla_{\nu} \phi(\varepsilon; \nu)] \\ &\approx \frac{1}{N_{MC}} \sum_{h=1}^{N_{MC}} \nabla_w \log p(Y|X, w)|_{w=\phi(\tilde{\varepsilon}^{(h)}; \nu)} \nabla_{\nu} \phi(\tilde{\varepsilon}^{(h)}; \nu), \quad \tilde{\varepsilon}^{(h)} \sim p(\varepsilon)\end{aligned}$$

Variational Inference - Reparameterization trick (Derivation)

Key observation

$$\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) = E_{p(\varepsilon)} \nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}$$

Now, we can turn the *gradient of an expectation* into an *expectation of a gradient*.

$$\begin{aligned}\nabla_{\nu} E_{q(w;\nu)} \log p(Y|X, w) &= \nabla_{\nu} E_{p(\varepsilon)} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \\ &= E_{p(\varepsilon)} [\nabla_{\nu} \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)}] \\ &= E_{p(\varepsilon)} [\nabla_w \log p(Y|X, w)|_{w=\phi(\varepsilon;\nu)} \nabla_{\nu} \phi(\varepsilon; \nu)] \\ &\approx \frac{1}{N_{MC}} \sum_{h=1}^{N_{MC}} \nabla_w \log p(Y|X, w)|_{w=\phi(\tilde{\varepsilon}^{(h)}; \nu)} \nabla_{\nu} \phi(\tilde{\varepsilon}^{(h)}; \nu), \quad \tilde{\varepsilon}^{(h)} \sim p(\varepsilon)\end{aligned}$$

Good news: if you use any auto-diff tool (PyTorch, Tensorflow, JAX, NumPyro, etc.), you will never compute this gradients manually.

Variational Inference - Reparameterization trick (Properties)

$$\nabla_{\boldsymbol{\nu}} \mathbb{E}_{q(\mathbf{w}; \boldsymbol{\nu})} \log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) \approx \frac{1}{N_{\text{MC}}} \sum_{h=1}^{N_{\text{MC}}} \nabla_{\mathbf{w}} \log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})|_{\mathbf{w}=\phi(\tilde{\boldsymbol{\varepsilon}}^{(h)}; \boldsymbol{\nu})} \nabla_{\boldsymbol{\nu}} \phi(\tilde{\boldsymbol{\varepsilon}}^{(h)}; \boldsymbol{\nu}), \quad \tilde{\boldsymbol{\varepsilon}}^{(h)} \sim p(\boldsymbol{\varepsilon})$$

- Estimation of the gradients is unbiased
- Need to be able to sample from $p(\boldsymbol{\varepsilon})$, but not from $q(\mathbf{w}; \boldsymbol{\nu})$
- The likelihood $p(\mathbf{Y}|\mathbf{X}, \mathbf{w})$ must be differentiable →
The neural network $f(\mathbf{x}; \mathbf{w})$ must be differentiable

Variational Inference with Stochastic Optimization

$$\widetilde{\mathcal{L}_{\text{ELBO}}} = \frac{1}{N_{\text{MC}}} \sum_{h=1}^{N_{\text{MC}}} \log p(\mathbf{Y}|\mathbf{X}, \tilde{\mathbf{w}}^{(h)}) - \text{KL}[q(\mathbf{w}; \nu) \| p(\mathbf{w})]$$

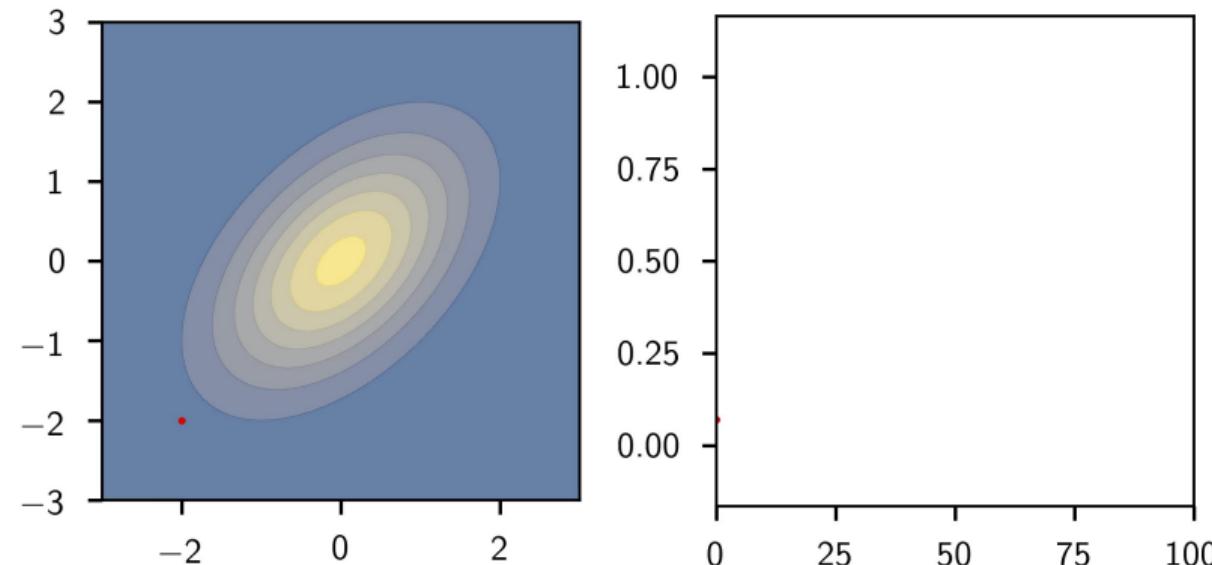
- We can get unbiased estimate by selecting M out of N data

$$\log p(\mathbf{Y}|\mathbf{X}, \tilde{\mathbf{w}}^{(h)}) \approx \frac{N}{M} \sum_{i \in \text{minibatch}} \log p(\mathbf{y}_i|\mathbf{x}_i, \tilde{\mathbf{w}}^{(h)})$$

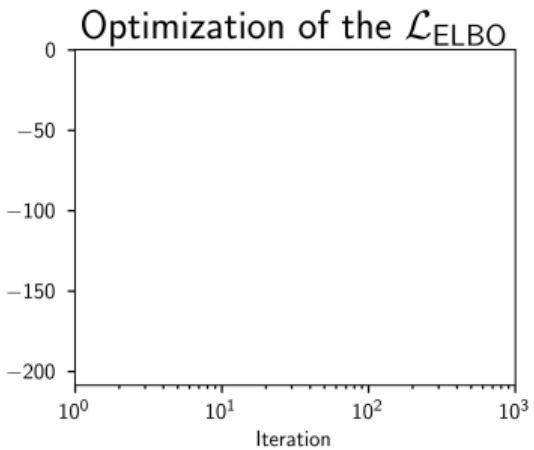
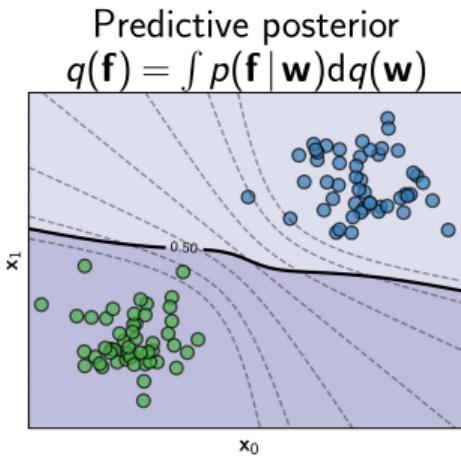
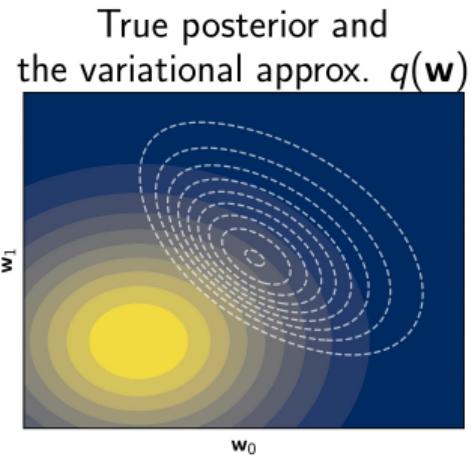
- We can use stochastic gradient optimization of our approximate variational objective with $\tilde{\mathbf{w}}^{(h)} = f(\tilde{\boldsymbol{\varepsilon}}^{(h)}; \nu)$ and $\tilde{\boldsymbol{\varepsilon}}^{(h)} \sim p(\boldsymbol{\varepsilon})$

Variational Inference with Stochastic Optimization

$$\nu' = \nu + \frac{\alpha_t}{2} \nabla_{\nu} (\widetilde{\mathcal{L}_{\text{ELBO}}}) \quad \alpha_t \rightarrow 0$$



A simple animation of stochastic variational inference in practice



Advances in Variational Inference for Bayesian Deep Learning

Variance analysis of Stochastic Variational Inference

Let $L_i = \log p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$ the likelihood contribution from $\{\mathbf{y}_i, \mathbf{x}_i\}$ in minibatch of size M

$$\text{Var} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] = N^2 \left(\frac{1}{M} \text{Var}[L_i] + \frac{M-1}{M} \text{Cov}[L_i, L_j] \right)$$

Variance analysis of Stochastic Variational Inference

Let $L_i = \log p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$ the likelihood contribution from $\{\mathbf{y}_i, \mathbf{x}_i\}$ in minibatch of size M

$$\text{Var} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] = N^2 \left(\frac{1}{M} \text{Var}[L_i] + \frac{M-1}{M} \text{Cov}[L_i, L_j] \right)$$

Variance analysis of Stochastic Variational Inference

Let $L_i = \log p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$ the likelihood contribution from $\{\mathbf{y}_i, \mathbf{x}_i\}$ in minibatch of size M

$$\text{Var} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] = N^2 \left(\frac{1}{M} \text{Var}[L_i] + \frac{M-1}{M} \text{Cov}[L_i, L_j] \right)$$

Solutions:

- Make $\text{Cov}[L_i, L_j] = 0$ by resampling $p(\epsilon)$ for every data-point (for each layer, $N_{MC} \times M \times D_{in} \times D_{out}$ times)
 - ▶ Computational intractable

Variance analysis of Stochastic Variational Inference

Let $L_i = \log p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$ the likelihood contribution from $\{\mathbf{y}_i, \mathbf{x}_i\}$ in minibatch of size M

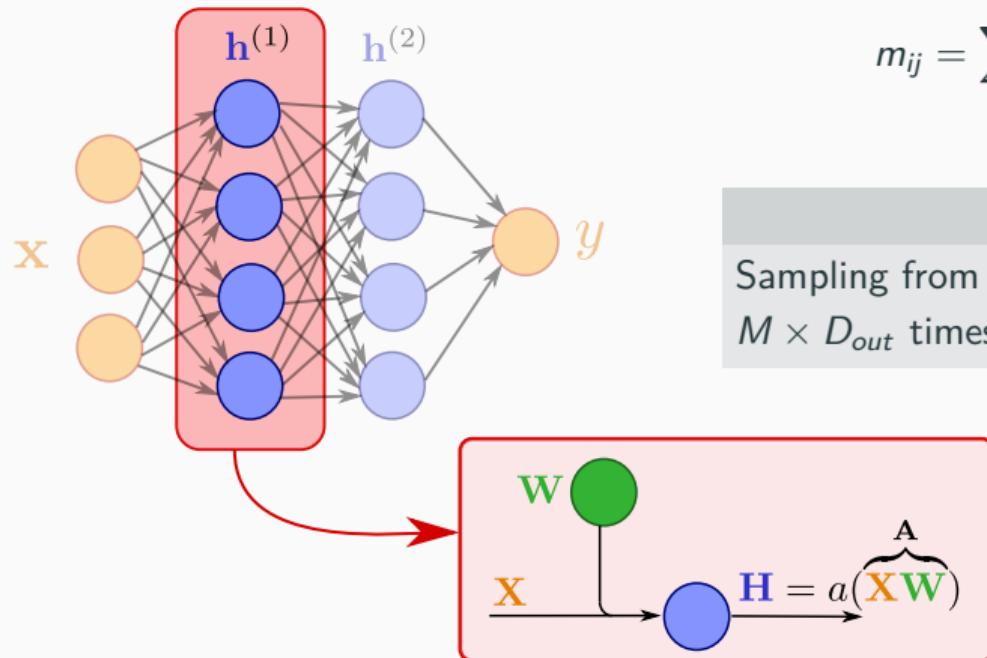
$$\text{Var} [\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] = N^2 \left(\frac{1}{M} \text{Var}[L_i] + \frac{M-1}{M} \text{Cov}[L_i, L_j] \right)$$

Solutions:

- Make $\text{Cov}[L_i, L_j] = 0$ by resampling $p(\epsilon)$ for every data-point (for each layer, $N_{MC} \times M \times D_{in} \times D_{out}$ times)
 - ▶ Computational intractable
- Move the stochasticity from the weights to the activations
 - ▶ The local reparameterization trick

The local reparameterization trick

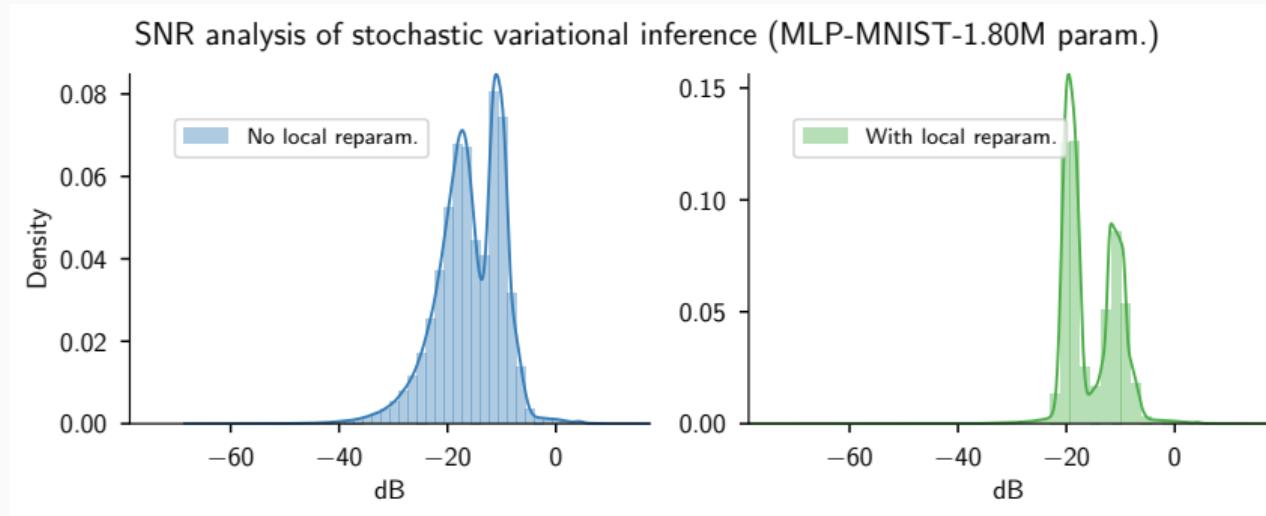
If $q(\mathbf{W}_{ij}) = \mathcal{N}(\mathbf{W}_{ij} | \boldsymbol{\mu}_{ij}, \sigma_{ij}^2)$,
then $q(A_{ij}) = \mathcal{N}(A_{ij} | m_{ij}, s_{ij}^2)$, with



$$m_{ij} = \sum_k \mathbf{X}_{ik} \boldsymbol{\mu}_{kj}, \quad s_{ij}^2 = \sum_k \mathbf{X}_{ik}^2 \sigma_{kj}^2$$

Sampling from $q(\mathbf{A})$ requires only to sample only $M \times D_{out}$ times from $q(\epsilon)$.

The local reparameterization trick (summary)



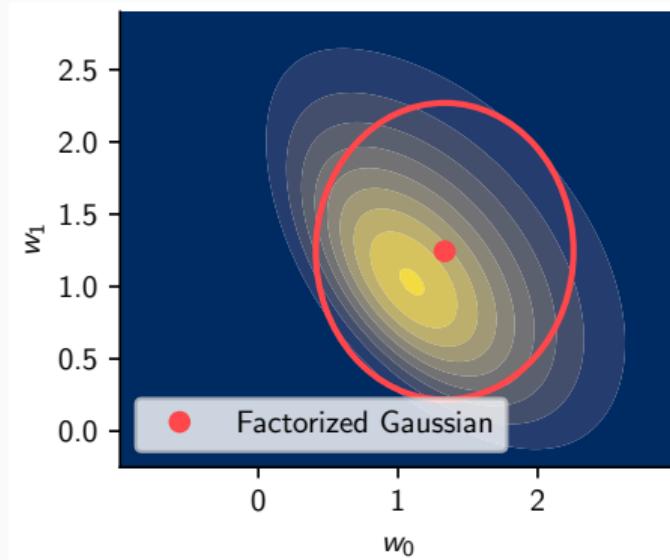
Practical improvements for training due to lower stochastic variance (speed-up convergence of $\mathcal{L}_{\text{ELBO}}$ or smaller batch-size)

Adapted from github.com/JavierAntoran/Bayesian-Neural-Networks

Richer family of parameterizations

Problem:

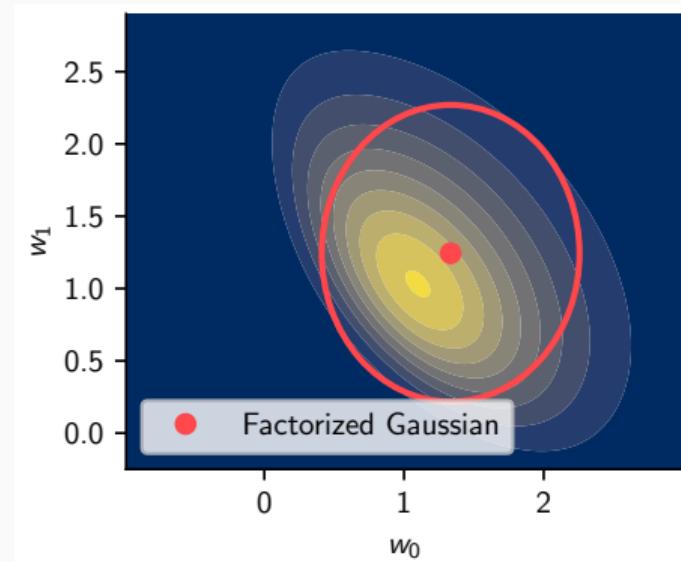
Mean-field Gaussians, albeit very simple to implement, is generally a rough approximation.



Richer family of parameterizations

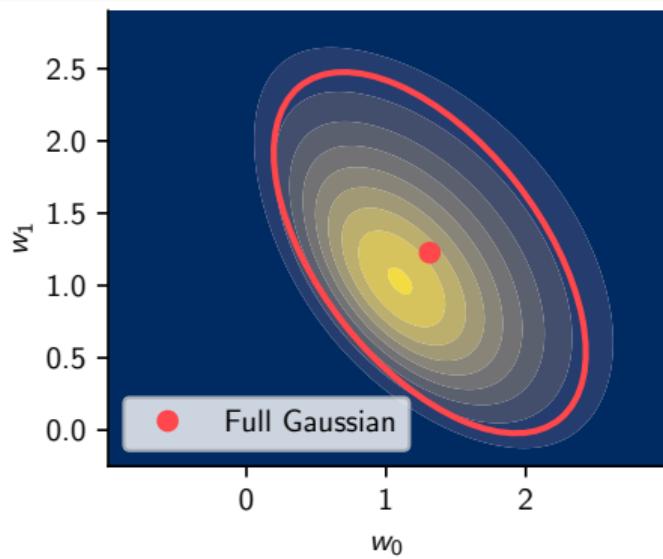
Problem:

Mean-field Gaussians, albeit very simple to implement, is generally a rough approximation.



Solution:

Introduce correlations among weights (i.e. Gaussian with non-diagonal covariance)



Here is where we deviate from standard VI methods

Working with full covariance is completely intractable

Assume we have one hidden layer with weights matrix \mathbf{W} of dimension 256×256 . From reshaping \mathbf{W} in vector form $\mathbf{w} \in \mathbb{R}^{65536}$, the covariance Σ is a matrix 65536×65536 .

Memory required: ~ 16 GB (only to store Σ , in single-point precision)

Here is where we deviate from standard VI methods

Working with full covariance is completely intractable

Assume we have one hidden layer with weights matrix \mathbf{W} of dimension 256×256 . From reshaping \mathbf{W} in vector form $\mathbf{w} \in \mathbb{R}^{65536}$, the covariance Σ is a matrix 65536×65536 .

Memory required: ~ 16 GB (only to store Σ , in single-point precision)

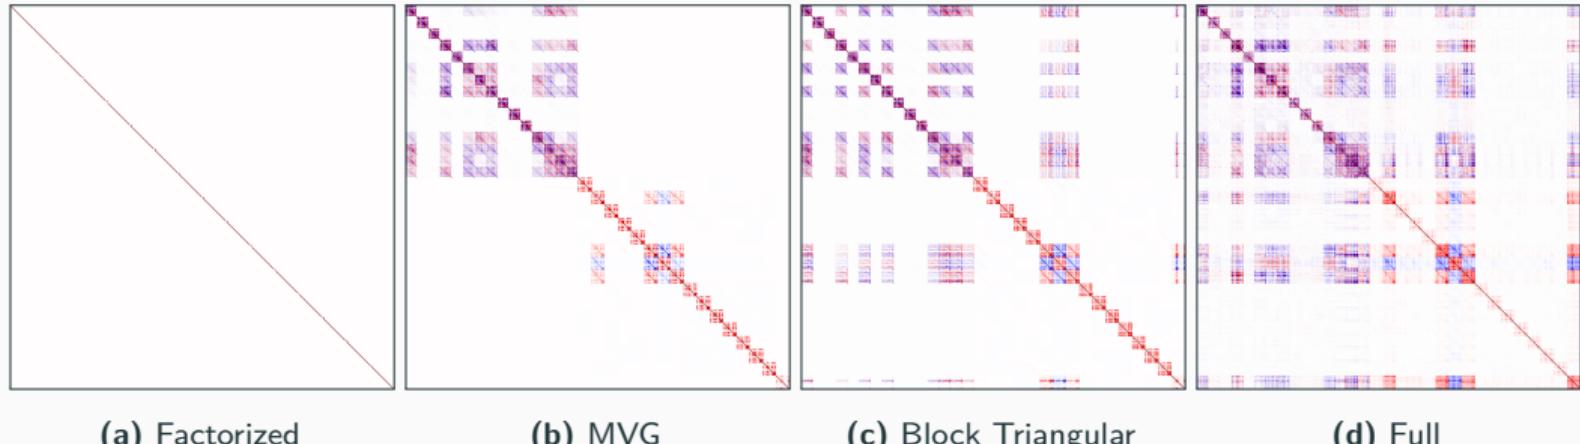
Simple solution:

Use mean-field for columns and full covariance for the rows of $\mathbf{W}_{:,i}$.

$$q(\mathbf{W}) = \prod_{i=1}^{256} q(\mathbf{W}_{:,i}) = \prod_{i=1}^{256} \mathcal{N}(\mathbf{W}_{:,i} | \boldsymbol{\mu}_i, \boldsymbol{\Sigma}_i)$$

where now $\boldsymbol{\Sigma}_i$ is a 256×256 covariance matrix.

Putting some structure in the covariance matrix



(a) Factorized

(b) MVG

(c) Block Triangular

(d) Full

Matrix Variate Gaussian distribution

$$q(\mathbf{W}) = \mathcal{MN}(\mathbf{W} | \mathbf{M}, \mathbf{U}, \mathbf{V}) = \frac{\exp\left(\frac{1}{2}\text{Tr}[\mathbf{V}^{-1}(\mathbf{W} - \mathbf{M})^\top \mathbf{U}^{-1}(\mathbf{W} - \mathbf{M})]\right)}{(2\pi)^{D_{in}D_{out}/2} |\mathbf{V}|^{D_{out}/2} |\mathbf{U}|^{D_{in}/2}}$$

$\mathbf{M} \in \mathbb{R}^{D_{in} \times D_{out}}$ is the mean, $\mathbf{U} \in \mathbb{R}^{D_{in} \times D_{in}}$ is the covariance matrix among rows and $\mathbf{V} \in \mathbb{R}^{D_{out} \times D_{out}}$ is the covariance matrix among columns.

Connected with Gaussian distribution:

$$\text{vec}(\mathbf{W}) \sim \mathcal{N}(\text{vec}(\mathbf{M}), \mathbf{V} \otimes \mathbf{U})$$

Matrix Variate Gaussian distribution

$$q(\mathbf{W}) = \mathcal{MN}(\mathbf{W} | \mathbf{M}, \mathbf{U}, \mathbf{V}) = \frac{\exp\left(\frac{1}{2}\text{Tr}[\mathbf{V}^{-1}(\mathbf{W} - \mathbf{M})^\top \mathbf{U}^{-1}(\mathbf{W} - \mathbf{M})]\right)}{(2\pi)^{D_{in}D_{out}/2} |\mathbf{V}|^{D_{out}/2} |\mathbf{U}|^{D_{in}/2}}$$

$\mathbf{M} \in \mathbb{R}^{D_{in} \times D_{out}}$ is the mean, $\mathbf{U} \in \mathbb{R}^{D_{in} \times D_{in}}$ is the covariance matrix among rows and $\mathbf{V} \in \mathbb{R}^{D_{out} \times D_{out}}$ is the covariance matrix among columns.

Reparameterization trick for $\mathbf{A} = \mathbf{X}\mathbf{W}$:

$$q(\mathbf{A}) = \mathcal{MN}(\mathbf{X}\mathbf{M}, \mathbf{X}\mathbf{U}\mathbf{X}^\top, \mathbf{V})$$

Structured posterior with connection to kernel methods

Impose a structure on the weight matrix inspired by scalable kernel methods

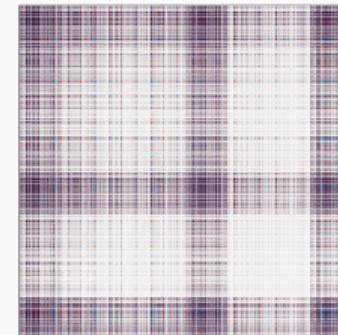
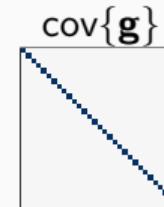
$$\textcolor{blue}{W} = \left(\begin{array}{c} \textcolor{violet}{S}_1 \\ \textcolor{purple}{W} \\ \textcolor{violet}{S}_2 \end{array} \right) \left(\begin{array}{c} \mathbf{H} \\ \mathbf{g} \sim \mathcal{N}(\mu, \Sigma) \\ \mathbf{H} \end{array} \right)$$

The diagram shows the factorization of the weight matrix $\textcolor{blue}{W}$. It is composed of three main components: $\textcolor{violet}{S}_1$ (top row), $\textcolor{purple}{W}$ (middle row), and $\textcolor{violet}{S}_2$ (bottom row). The middle component $\textcolor{purple}{W}$ is further factored into \mathbf{H} , $\mathbf{g} \sim \mathcal{N}(\mu, \Sigma)$, and \mathbf{H} . The matrices \mathbf{H} and \mathbf{g} are shown as 4x4 grids. The grid for \mathbf{H} contains values +1 and -1. The grid for \mathbf{g} is a 4x4 identity matrix. The $\textcolor{violet}{S}_1$ and $\textcolor{violet}{S}_2$ components are diagonal matrices with purple blocks.

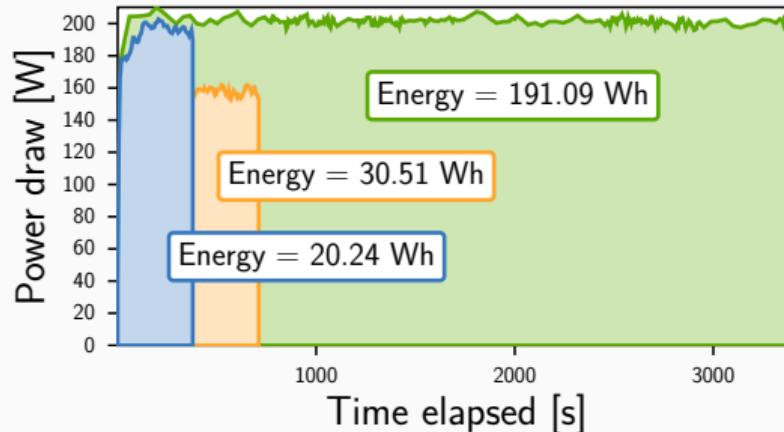
Benefit:

- Space complexity for storing $\textcolor{blue}{W}$:
 $\mathcal{O}(D^2) \rightarrow \mathcal{O}(D)$.
- Time complexity for $\textcolor{blue}{W}\mathbf{x}$:
 $\mathcal{O}(D^2) \rightarrow \mathcal{O}(D \log D)$.

$$\text{cov}\{\text{vect}(\textcolor{blue}{W})\}$$



Keeping an eye to energy efficiency



Comparison between Monte Carlo dropout (●), Matrix Variate Gaussian (●) and Hadamard factorization (●).

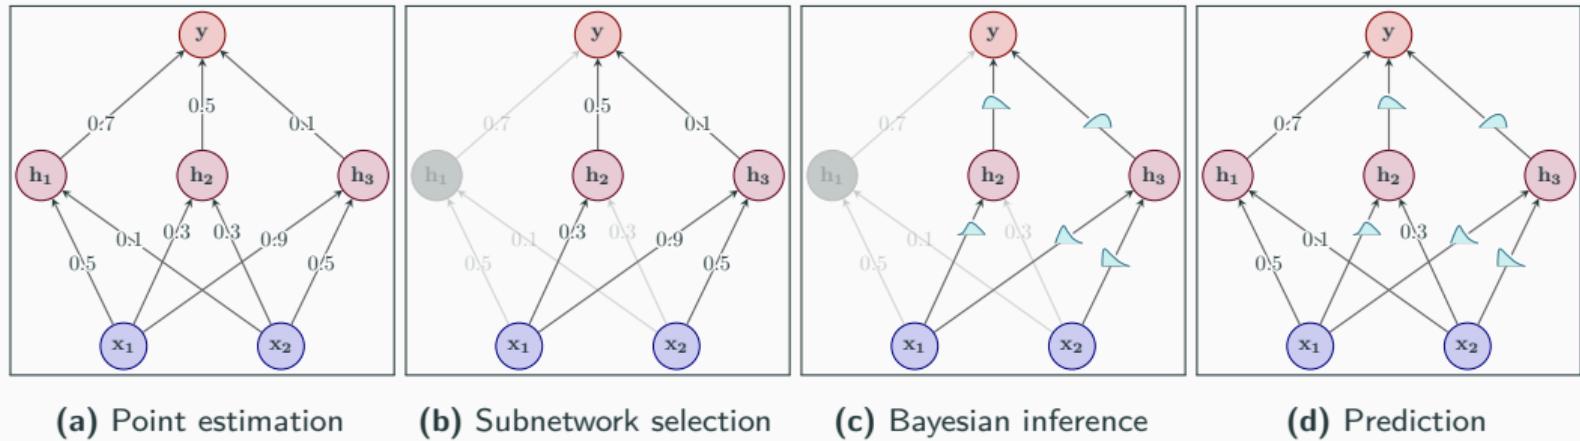
Low rank factorization with variational inference

Performing tensor decomposition with variational inference is not straightforward.



The redundant variational parameterization induced by the tensor cores makes the optimization landscapes highly multi-modal, thus leading to slow convergence.

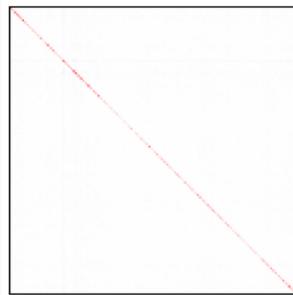
Imposing low-rank structure by doing inference on a subset of weights



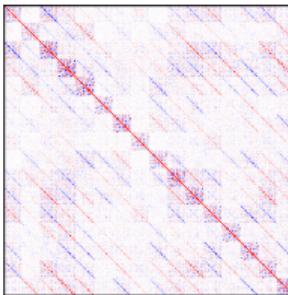
- (a) Train a neural network to a MAP solution
- (b) Identify a small subset of the weights
- (c) Estimate a posterior distribution over the selected subset
- (d) Predict using the mix of Bayesian and deterministic weights

Complex covariances can raise from mean-field and depth

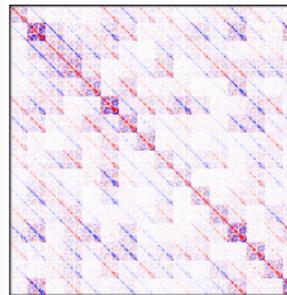
Covariance heatmap for mean-field approximate posteriors trained on FashionMNIST.



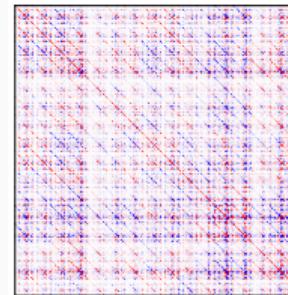
(a) One weight
matrix.



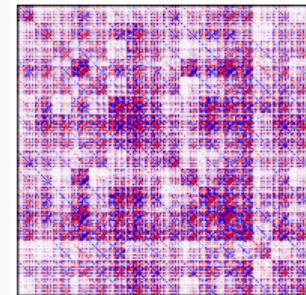
(b) 5-layers (linear)



(c) 10-layers (linear)



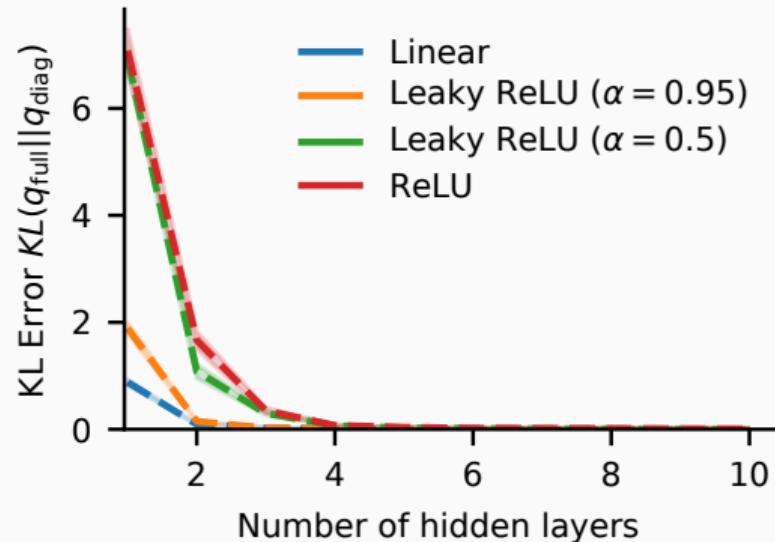
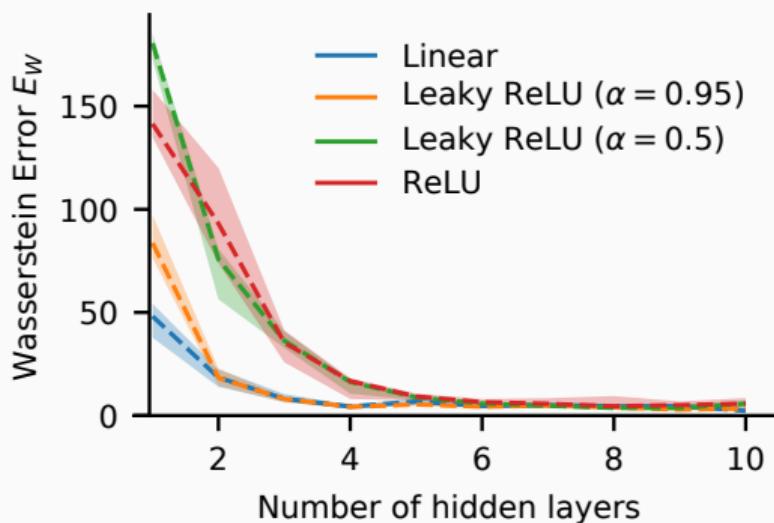
(d) 5-layers (Leaky
ReLU)



(e) 10-layers (Leaky
ReLU)

The tradeoff depth/structure

For any sufficiently deep and wide neural network, there exists a mean-field distribution which induces the same distribution over function values as that induced by the posterior predictive



Farquhar et al. (2020). *Liberty or Depth: Deep Bayesian Neural Nets Do Not Need Complex Weight Posterior Approximations*. NeurIPS

Are structured posteriors worth? (Use case: CIFAR10)

For low-medium depth models having structured posterior seems to be important.

| Architecture | Covariance | Accuracy (\uparrow) | NLL (\downarrow) | ECE (\downarrow) |
|-----------------------------|-----------------|-------------------------|----------------------|----------------------|
| AlexNet (low depth) | Diagonal | 75.5% | 0.703 | 0.016 |
| | Low-rank (WHVI) | 88.5% | 0.490 | 0.009 |
| ResNet-18 (medium depth) | Diagonal | 84.3% | 0.477 | 0.040 |
| | Low-rank (WHVI) | 86.4% | 0.616 | 0.029 |

Farquhar et al. (2020). *Liberty or Depth: Deep Bayesian Neural Nets Do Not Need Complex Weight Posterior Approximations*. NeurIPS

Rossi et al. (2020). *Walsh-Hadamard Variational Inference for Bayesian Deep Learning*. NeurIPS

Osawa et al. (2019). *Practical Deep Learning with Bayesian Principles*. NeurIPS

Are structured posteriors worth? (Use case: ImageNet)

The importance of structured covariance seems to be diminished in very large-scale models.

| Architecture | Covariance | Accuracy (\uparrow) | NLL (\downarrow) | ECE (\downarrow) |
|--------------|------------|-------------------------|----------------------|----------------------|
| DenseNet-161 | Diagonal | 78.6% | 0.86 | 0.046 |
| | Low-rank | 78.6% | 0.83 | 0.020 |
| ResNet-152 | Diagonal | 80.0% | 0.86 | 0.057 |
| | Low-rank | 79.1% | 0.82 | 0.028 |

Farquhar et al. (2020). *Liberty or Depth: Deep Bayesian Neural Nets Do Not Need Complex Weight Posterior Approximations*. NeurIPS

Introduction to Variational Inference

- Jordan et al. (1999). *An Introduction to Variational Methods for Graphical Models*. Mach. Learn.
- Hoffman et al. (2013). *Stochastic Variational Inference*. JMLR
- Ranganath et al. (2014). *Black Box Variational Inference*. AISTATS
- Blei et al. (2017). *Variational Inference: A Review for Statisticians*. JASA

Monte-Carlo Dropout for Bayesian Neural Networks and follow-up

- Srivastava et al. (2014). *Dropout: A Simple Way to Prevent Neural Networks from Overfitting*. JMLR
- Kingma et al. (2015). *Variational Dropout and the Local Reparameterization Trick*. NeurIPS
- Gal (2016). *Uncertainty in Deep Learning*. University of Cambridge (PhD Thesis)
- Gal and Ghahramani (2016). *Bayesian Convolutional Neural Networks with Bernoulli Approximate Variational Inference*. ICLR Workshop
- Gal and Ghahramani (2016). *Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning*. ICML
- Kendall and Gal (2017). *What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?*. NeurIPS
- Li and Gal (2017). *Dropout Inference in Bayesian Neural Networks with Alpha-divergences*. ICML

- Hron et al. (2017). *Variational Gaussian Dropout is not Bayesian*. NeurIPS Workshop
- Hron et al (2018). *Variational Bayesian Dropout: Pitfalls and Fixes*. ICML

Variational Inference for Neural Networks

- Graves (2011). *Practical Variational Inference for Neural Networks*. NeurIPS
- Rezende et al. (2014). *Stochastic Backpropagation and Approximate Inference in Deep Generative Models*. ICML
- Hernández-Lobato et al. (2015). *Probabilistic Backpropagation for Scalable Learning of Bayesian Neural Networks*. ICML
- Blundell et al. (2015). *Weight Uncertainty in Neural Networks*. ICML
- Rezende et al. (2015). *Variational Inference with Normalizing Flows*. ICML
- Louizos and Welling (2016). *Structured and Efficient Variational Deep Learning with Matrix Gaussian Posteriors*. ICML
- Kingma et al. (2016). *Improving Variational Inference with Inverse Autoregressive Flow*. NeurIPS

References v

- Liu et al. (2016). *Stein Variational Gradient Descent: A General Purpose Bayesian Inference Algorithm*. NeurIPS
- Miller et al. (2016). *Variational Boosting: Iteratively Refining Posterior Approximations*. ICML
- Louizos and Welling (2017). *Multiplicative Normalizing Flows for Variational Bayesian Neural Networks*. ICML
- Sun et al. (2017). *Learning Structured Weight Uncertainty in Bayesian Neural Networks*. AISTATS
- Khan et al. (2018). *Fast and Scalable Bayesian Deep Learning by Weight-Perturbation in ADAM*. ICML
- Rossi et al. (2018). *Good Initializations of Variational Bayes for Deep Models*. ICML
- Zhang et al. (2018). *Noisy Natural Gradient as Variational Inference*. ICML
- Ghosh et al. (2018). *Structured Variational Learning of Bayesian Neural Networks with Horseshoe Priors*. ICML
- Osawa et al. (2019). *Practical Deep Learning with Bayesian Principles*. NeurIPS

- Sun et al. (2019). *Functional Variational Bayesian Neural Networks*. ICLR
- Farquhar et al. (2020). *Liberty or Depth: Deep Bayesian Neural Nets Do Not Need Complex Weight Posterior Approximations*. NeurIPS
- Rossi et al. (2020). *Walsh-Hadamard Variational Inference for Bayesian Deep Learning*. NeurIPS
- Daxberger et al. (2021). *Bayesian Deep Learning via Subnetwork Inference*. ICML